



Snowflake QA Analytics Bot

This project demonstrates how to build and deploy a **Question–Answer (QA) analytics bot** using **Snowflake Cortex**, **Snowpark**, and **Streamlit**.

Users can ask questions in **natural language**, and the system answers them by running **SQL queries on Snowflake data**.

The application is:

- Publicly accessible
 - Secure
 - Read-only
 - Does not expose Snowflake credentials
-

📌 What I Built

I built a **web-based QA bot** that:

- Accepts questions in English
 - Uses **Snowflake Cortex** to convert questions into SQL
 - Executes SQL queries on Snowflake data
 - Displays results in a Streamlit UI
 - Is deployed publicly using **Render**
-

🏗 Architecture Overview



🗄 Snowflake Data Used

Database

DEMO_DB

Schema

PUBLIC

Tables

CUSTOMERS

- CUSTOMER_ID
- CUSTOMER_NAME
- CITY

SALES

- ORDER_ID
- CUSTOMER_ID
- AMOUNT
- ORDER_DATE

The relationship between tables:

SALES.CUSTOMER_ID = CUSTOMERS.CUSTOMER_ID

This allows analytics such as:

- Total sales per customer
- Orders by customer
- Sales trends

Warehouse Used

Warehouse

DEMO_WH

This warehouse:

- Executes all queries
- Is cost-efficient
- Automatically starts and stops

- Is suitable for demo and learning workloads
-



Semantic View (Understanding Layer)

A **Semantic View (YAML)** was created to describe:

- Business meaning of columns
- Relationships between tables
- Common analytical questions

Important:

- Semantic Views **do NOT store data**
 - Semantic Views **cannot be queried directly**
 - They are used only to help **Cortex understand the data**
-



Execution SQL View (Why & How)

Why this view was needed

Snowflake does not allow:

```
SELECT * FROM SEMANTIC_VIEW
```

So a **normal SQL VIEW** was created for execution.

View Created thru semantic view yaml file.

```
CREATE OR REPLACE VIEW DEMO_DB.PUBLIC.SHIVANSHI_EXEC_VIEW AS
SELECT
    c.CUSTOMER_ID,
    c.CUSTOMER_NAME,
    s.ORDER_ID,
    s.ORDER_DATE,
    s.AMOUNT
FROM DEMO_DB.PUBLIC.SALES s
INNER JOIN DEMO_DB.PUBLIC.CUSTOMERS c
    ON s.CUSTOMER_ID = c.CUSTOMER_ID;
```

This view:

- Does NOT duplicate data
- Is safe for large datasets

- Provides a single execution surface for AI-generated SQL
-

Role & Security Setup

A **read-only role** was created following least privilege security.

Role Name

SPCS_AGENT_ROLE

Permissions Granted

- Read access to tables and views
- Usage access to warehouse
- Access to Snowflake Cortex
- **✗ No write or admin permissions**

This role is assigned to the Snowflake user used by the app.

Streamlit Application (`app.py`)

The Streamlit app:

- Displays a UI
 - Takes user questions
 - Enforces a minimum of 4 words
 - Disables Snowflake query caching
 - Uses Cortex to generate SQL
 - Executes SQL on the execution view
 - Displays results in a table
-

Full `app.py` Code

```
import streamlit as st
from snowflake.snowpark import Session
import os

# -----
```

```

# Page Config
# -----
st.set_page_config(
    page_title="QA Bot",
    page_icon="🤖",
    layout="centered"
)

# -----
# Header
# -----
st.markdown(
    """
<h1 style="text-align:center;">🤖 QA Bot</h1>
<p style="text-align:center; color: gray;">
    QA bot powered by Snowflake Cortex & Semantic View
</p>
"""
,
unsafe_allow_html=True
)

st.divider()

# -----
# Example Questions
# -----
st.markdown("###💡 Try an example question")

example_questions = [
    "List all customers and their orders",
    "Which customer has the highest total sales",
    "Show total sales by city",
]

cols = st.columns(len(example_questions))
for i, q in enumerate(example_questions):
    if cols[i].button(q):
        st.session_state["question"] = q

st.divider()

# -----
# Question Input
# -----
question = st.text_input(
    "Ask a question about Sales or Customers",
    key="question",
    placeholder="e.g. Which customer has the highest total sales"
)

# -----
# Ask Button
# -----
if st.button("Ask 🤖", use_container_width=True):

    # Rule: Minimum 4 words

```

```

if len(question.strip().split()) < 4:
    st.error("X Invalid Length")
    st.stop()

with st.spinner("Thinking..."):

    # Snowflake connection
    session = Session.builder.configs({
        "account": os.getenv("SNOWFLAKE_ACCOUNT"),
        "user": os.getenv("SNOWFLAKE_USER"),
        "password": os.getenv("SNOWFLAKE_PASSWORD"),
        "role": os.getenv("SNOWFLAKE_ROLE"),
        "warehouse": "DEMO_WH",
        "database": "DEMO_DB",
        "schema": "PUBLIC"
    }).create()

    # Disable caching
    session.sql("ALTER SESSION SET USE_CACHED_RESULT = FALSE").collect()

    # Cortex prompt
    prompt = f"""
You are a Snowflake Cortex Analyst.

You understand the business using this semantic model:
- Customers place sales orders
- SALES.CUSTOMER_ID joins to CUSTOMERS.CUSTOMER_ID
- Measures: AMOUNT
- Time dimension: ORDER_DATE

Generate ONE valid Snowflake SELECT query.
Do not explain.
Do not include markdown.
Do not include comments.

Execute queries ONLY on this view:
DEMO_DB.PUBLIC.SHIVANSHI_EXEC_VIEW

Question:
{question}
"""

    cortex_sql = f"""
SELECT SNOWFLAKE.CORTEX.COMPLETE(
    'snowflake-arctic',
    $$ {prompt} $$
) AS SQL_QUERY
"""

    generated_sql = session.sql(cortex_sql).collect()[0]["SQL_QUERY"]

    if not generated_sql.strip().lower().startswith("select"):
        st.error("⚠️ Unable to generate a valid query")
        st.stop()

    df = session.sql(generated_sql).to_pandas()

```

```
session.close()

st.success("✅ Answer generated")

with st.expander("🔍 View generated SQL"):
    st.code(generated_sql, language="sql")

st.markdown("### 📊 Result")
st.dataframe(df, use_container_width=True)
```

🔑 Credential Management

Snowflake credentials are stored as **environment variables** (in Render):

- SNOWFLAKE_ACCOUNT
- SNOWFLAKE_USER
- SNOWFLAKE_PASSWORD
- SNOWFLAKE_ROLE

They are never stored in code and never visible to users.

🚀 Deployment

- Code pushed to GitHub
- Streamlit app deployed on Render
- Environment variables configured in Render
- Public URL generated

Users can access the app without logging into Snowflake.